

FPGA Implementation of the Nintendo Entertainment System (NES)

Four People Generating A Nintendo Entertainment System (FPGANES)

Eric Sullivan, Jonathan Ebert, Patrick Yang, Pavan Holla

Final Report

University of Wisconsin-Madison

ECE 554

Spring 2017

Introduction	6
Top Level Block Diagram	7
Top level description	7
Data Flow Diagram	8
Control Flow Diagram	8
CPU	9
CPU Registers	9
CPU ISA	9
CPU Addressing Modes	10
CPU Interrupts	10
CPU Opcode Matrix	10
CPU Block Diagram	13
CPU Top Level Interface	15
CPU Instruction Decode Interface	15
CPU MEM Interface	16
CPU ALU Interface	16
ALU Input Selector	17
CPU Registers Interface	17
CPU Processor Control Interface	18
CPU Enums	19
Picture Processing Unit	20
PPU Top Level Schematic	21
PPU Memory Map	22
PPU CHAROM	22
PPU Rendering	23
PPU Memory Mapped Registers	25

PPU Register Block Diagram	26
PPU Register Descriptions	26
PPU Background Renderer	28
PPU Background Renderer Diagram	29
PPU Sprite Renderer	30
PPU Sprite Renderer Diagram	32
PPU Object Attribute Memory	33
PPU Palette Memory	33
VRAM Interface	34
DMA	34
PPU Testbench	35
PPU Testbench PPM file format	36
PPU Testbench Example Renderings	37
Memory Maps	37
PPU ROM Memory Map	37
CPU ROM Memory Map	38
Memory Mappers Interface	39
APU	40
APU Registers	40
Controllers (SPART)	42
Debug Modification	42
Controller Registers	42
Controllers Wrapper	42
Controller Wrapper Diagram	43
Controller Wrapper Interface	43
Controller	44
Controller Diagram	44

Controller Interface	44
Special Purpose Asynchronous Receiver and Transmitter (SPART)	45
SPART Diagram & Interface	45
Controller Driver	46
Controller Driver State Machine	46
VGA	47
VGA Diagram	47
VGA Interface	48
VGA Clock Gen	48
VGA Timing Gen	49
VGA Display Plane	49
VGA RAM Wrapper	50
VGA RAM Reader	50
Software	52
Controller Simulator	52
Controller Simulator State Machine	52
Controller Simulator Output Packet Format	52
Controller Simulator GUI and Button Map	53
Assembler	53
Opcode Table	54
NES Assembly Formats	56
Invoking Assembler	57
iNES ROM Converter	57
Tic Tac Toe	57
Testing & Debug	58
Simulation	58
Test	58

Integrated Simulation	58
Tracer	58
Results	59
Possible Improvements	59
References and Links	59
Contributions	60
Eric Sullivan	60
Patrick Yang	60
Pavan Holla	60
Jonathan Ebert	60

1. Introduction

Following the video game crash in the early 1980s, Nintendo released their first video game console, the Nintendo Entertainment System (NES). Following a slow release and early recalls, the console began to gain momentum in a market that many thought had died out, and the NES is still appreciated by enthusiasts today. A majority of its early success was due to the relationship that Nintendo created with third-party software developers. Nintendo required that restricted developers from publishing games without a license distributed by Nintendo. This decision led to higher quality games and helped to sway the public opinion on video games, which had been plagued by poor games for other gaming consoles.

Our motivation is to better understand how the NES worked from a hardware perspective, as the NES was an extremely advanced console when it was released in 1985 (USA). The NES has been recreated multiple times in software emulators, but has rarely been done in a hardware design language, which makes this a unique project. Nintendo chose to use the 6502 processor, also used by Apple in the Apple II, and chose to include a picture processing unit to provide a memory efficient way to output video to the TV. Our goal was to recreate the CPU and PPU in hardware, so that we could run games that were run on the original console. In order to exactly recreate the original console, we needed to include memory mappers, an audio processing unit, a DMA unit, a VGA interface, and a way to use a controller for input. In addition, we wrote our own assembler and tic-tac-toe game to test our implementation. The following sections will explain the microarchitecture of the NES. Much of the information was gleaned from nesdev.com, and from other online forums that reverse engineered the NES.

2. Top Level Block Diagram

2.1.Top level description

Here is an overview of each module in our design. Our report has a section dedicated for each of these modules.

2.1.1.PPU - The PPU(Picture Processing Unit) is responsible for rendering graphics on screen. It receives memory mapped accesses from the CPU, and renders graphics from memory, providing RGB values.

2.1.2.CPU - Our CPU is a 6502 implementation. It is responsible for controlling all other modules in the NES. At boot, CPU starts reading programs at the address 0xFFFC.

2.1.3.DMA - The DMA transfers chunks of data from CPU address space to PPU address space. It is faster than performing repeated Loads and Stores in the CPU.

2.1.4.Display Memory and VGA - The PPU writes to the display memory, which is subsequently read out by the VGA module. The VGA module produces the hsync, vsync and RGB values that a monitor requires.

2.1.5.Controller - A program runs on a host computer which transfers serial data to the FPGA. The protocol used by the controller is UART in our case

2.1.6.APU - Generates audio in the NES. However, we did not implement this module.

2.1.7.CHAR RAM/ RAM - Used by the CPU and PPU to store temporary data

2.1.8.PROG ROM/ CHAR ROM - PROG ROM contains the software(instructions) that runs the game. CHAR ROM on the other hand contains mostly image data and graphics used in the game.

2.2.Data Flow Diagram

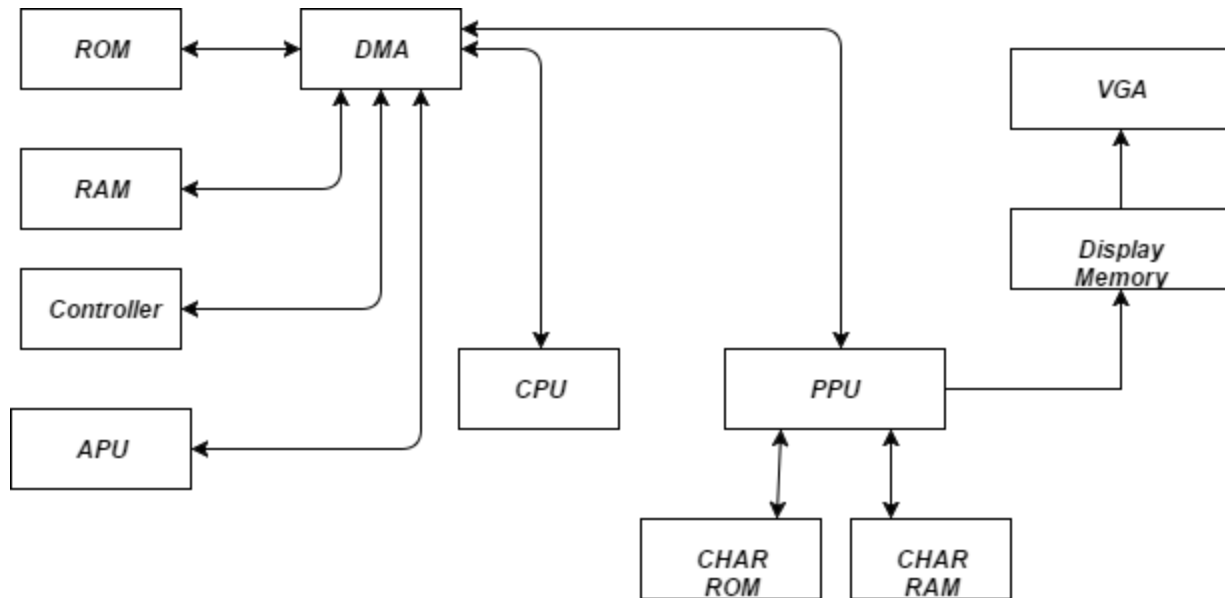


Figure 1: System level data flow diagram

2.3. Control Flow Diagram

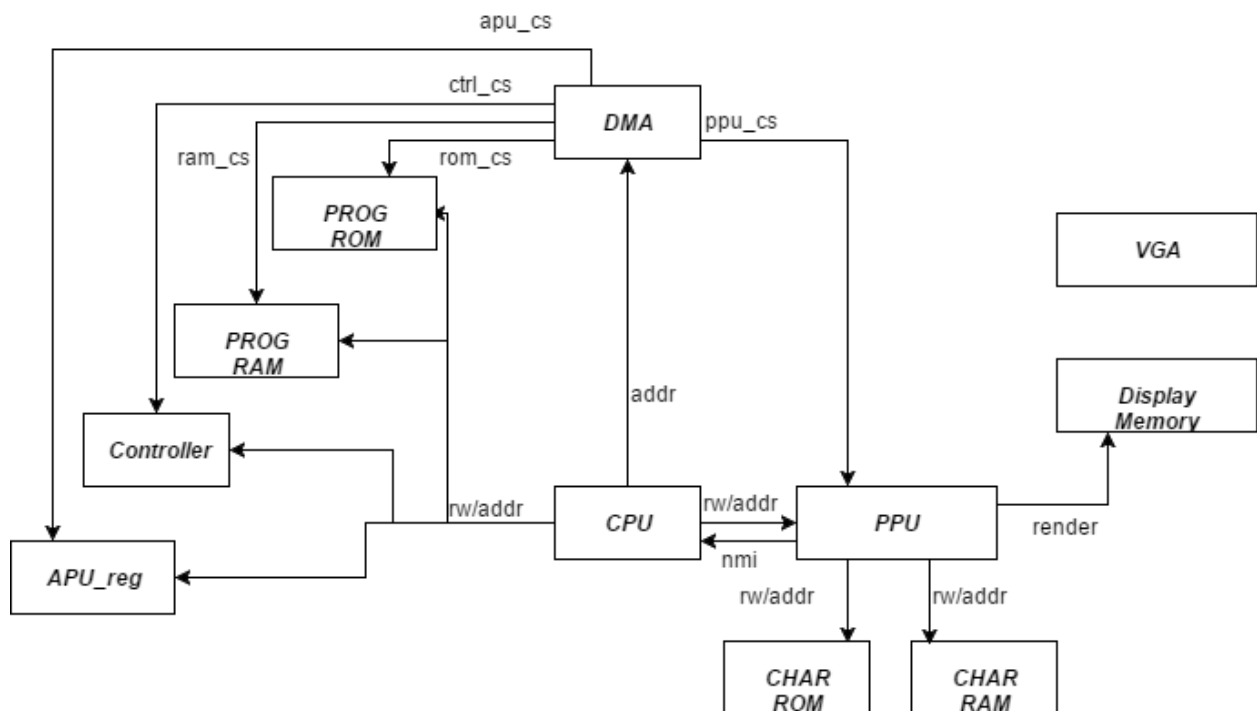


Figure 2: System level control flow diagram.

3. CPU

CPU Registers

The CPU of the NES is the MOS 6502. It is an accumulator plus index register machine. There are five primary registers on which operations are performed:

1. PC : Program Counter
2. Accumulator(A) : Keeps track of results from ALU
3. X : X index register
4. Y : Y index register
5. Stack pointer
6. Status Register : Negative, Overflow, Unused, Break, Decimal, Interrupt, Zero, Carry
 - Break means that the current interrupt is from software interrupt, BRK
 - Interrupt is high when maskable interrupts (IRQ) is to be ignored. Non-maskable interrupts (NMI) cannot be ignored.

There are 6 secondary registers:

1. AD : Address Register
 - Stores where to jump to or where to get indirect access from.
2. ADV : AD Value Register
 - Stores the value from indirect access by AD.
3. BA : Base Address Register
 - Stores the base address before index calculation. After the calculation, the value is transferred to AD if needed.
4. BAV : BA Value Register
 - Stores the value from indirect access by BA.
5. IMM : Immediate Register
 - Stores the immediate value from the memory.
6. Offset
 - Stores the offset value of branch from memory

CPU ISA

The ISA may be classified into a few broad operations:

- Load into A,X,Y registers from memory
- Perform arithmetic operation on A,X or Y
- Move data from one register to another
- Program control instructions like Jump and Branch
- Stack operations
- Complex instructions that read, modify and write back memory.

CPU Addressing Modes

Additionally, there are thirteen addressing modes which these operations can use. They are

- **Accumulator** – The data in the accumulator is used.
- **Immediate** - The byte in memory immediately following the instruction is used.
- **Zero Page** – The Nth byte in the first page of RAM is used where N is the byte in memory immediately following the instruction.
- **Zero Page, X Index** – The (N+X)th byte in the first page of RAM is used where N is the byte in memory immediately following the instruction and X is the contents of the X index register.
- **Zero Page, Y Index** – Same as above but with the Y index register
- **Absolute** – The two bytes in memory following the instruction specify the absolute address of the byte of data to be used.
- **Absolute, X Index** - The two bytes in memory following the instruction specify the base address. The contents of the X index register are then added to the base address to obtain the address of the byte of data to be used.
- **Absolute, Y Index** – Same as above but with the Y index register
- **Implied** – Data is either not needed or the location of the data is implied by the instruction.
- **Relative** – The content of sum of (the program counter and the byte in memory immediately following the instruction) is used.
- **Absolute Indirect** - The two bytes in memory following the instruction specify the absolute address of the two bytes that contain the absolute address of the byte of data to be used.
- **(Indirect, X)** – A combination of Indirect Addressing and Indexed Addressing
- **(Indirect), Y** - A combination of Indirect Addressing and Indexed Addressing

CPU Interrupts

The 6502 supports three interrupts. The reset interrupt routine is called after a physical reset. The other two interrupts are the non_maskable_interrupt(NMI) and the general_interrupt(IRQ). The general_interrupt can be disabled by software whereas the others cannot. When interrupt occurs, the CPU finishes the current instruction then PC jumps to the specified interrupt vector then return when finished.

CPU Opcode Matrix

The NES 6502 ISA is a CISC like ISA with 56 instructions. These 56 instructions can pair up with addressing modes to form various opcodes. The opcode is always 8 bits, however based on the addressing mode, upto 4 more memory location may need to be fetched. The memory is single cycle, i.e data[7:0] can be latched the cycle after address[15:0] is placed on the bus. The following tables summarize the instructions available and possible addressing modes:

Storage	
LDA	Load A with M
LDX	Load X with M
LDY	Load Y with M
STA	Store A in M
STX	Store X in M

STY	Store Y in M
TAX	Transfer A to X
TAY	Transfer A to Y
TSX	Transfer Stack Pointer to X
TXA	Transfer X to A
TXS	Transfer X to Stack Pointer
TYA	Transfer Y to A
Arithmetic	
ADC	Add M to A with Carry
DEC	Decrement M by One
DEX	Decrement X by One
DEY	Decrement Y by One
INC	Increment M by One
INX	Increment X by One
INY	Increment Y by One
SBC	Subtract M from A with Borrow
Bitwise	
AND	AND M with A
ASL	Shift Left One Bit (M or A)
BIT	Test Bits in M with A
EOR	Exclusive-Or M with A
LSR	Shift Right One Bit (M or A)
ORA	OR M with A
ROL	Rotate One Bit Left (M or A)
ROR	Rotate One Bit Right (M or A)
Branch	
BCC	Branch on Carry Clear
BCS	Branch on Carry Set
BEQ	Branch on Result Zero
BMI	Branch on Result Minus
BNE	Branch on Result not Zero
BPL	Branch on Result Plus
BVC	Branch on Overflow Clear
BVS	Branch on Overflow Set
Jump	
JMP	Jump to Location
JSR	Jump to Location Save Return Address
RTI	Return from Interrupt
RTS	Return from Subroutine
Status Flags	

CLC	Clear Carry Flag
CLD	Clear Decimal Mode
CLI	Clear interrupt Disable Bit
CLV	Clear Overflow Flag
CMP	Compare M and A
CPX	Compare M and X
CPY	Compare M and Y
SEC	Set Carry Flag
SED	Set Decimal Mode
SEI	Set Interrupt Disable Status
Stack	
PHA	Push A on Stack
PHP	Push Processor Status on Stack
PLA	Pull A from Stack
PLP	Pull Processor Status from Stack
System	
BRK	Force Break
NOP	No Operation

The specific opcode hex values are specified in the Assembler section.

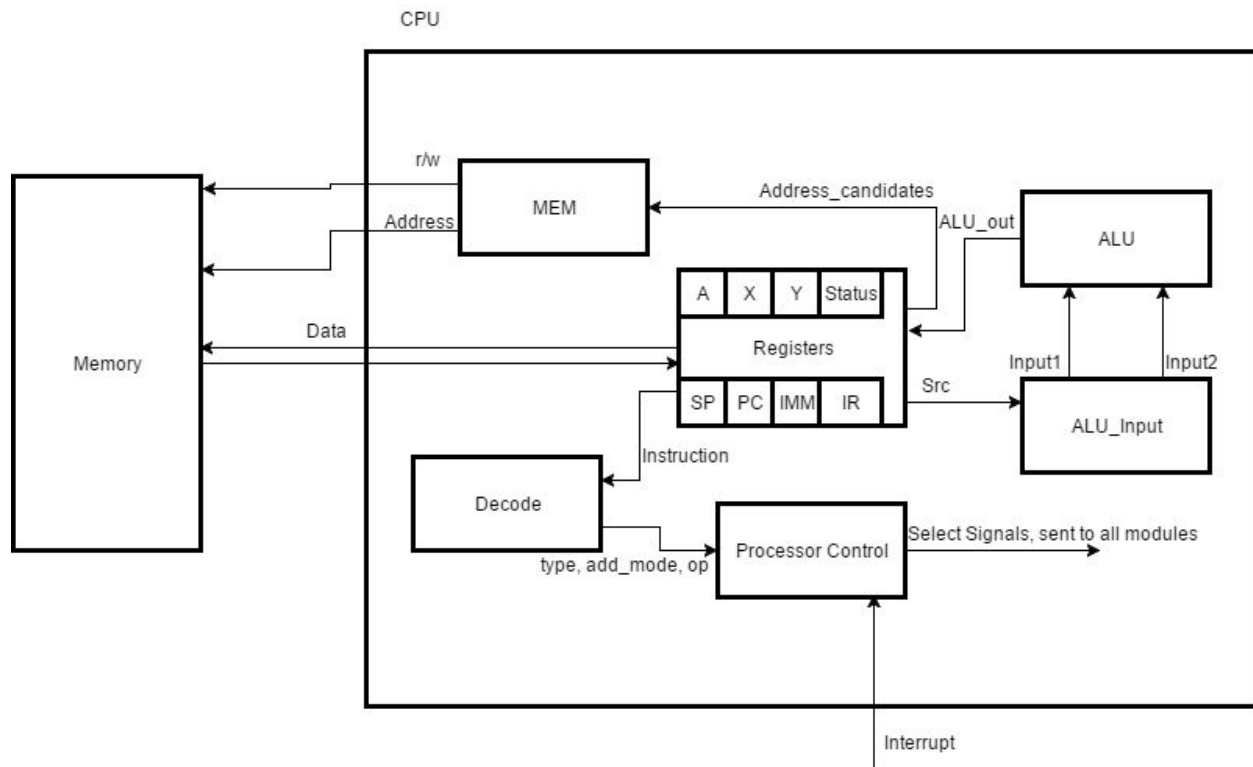
For more information on the opcodes, please refer

<http://www.6502.org/tutorials/6502opcodes.html>

or

http://www.thealmightyguru.com/Games/Hacking/Wiki/index.php/6502_Opcodes

CPU Block Diagram



Block	Primary Function
Decode	Decode the current instruction. Classifies the opcode into an instruction_type(arithmetic,ld etc) and addressing mode(immediate, indirect etc)
Processor Control	State machine that keeps track of current instruction stage, and generates signals to load registers.
ALU	Performs ALU ops and handles Status Flags
Registers	Contains all registers. Register values change according to signals from processor control.
Mem	Acts as the interface between CPU and memory. Mem block thinks it's communicating with the memory but the DMA can reroute the communication to any other blocks like PPU, controller

Instruction flow

The following table presents a high level overview of how each instruction is handled.

Cycle Number	Blocks	Action
0	Processor Control → Registers	Instruction Fetch
1	Register → Decode	Classify instruction and addressing mode
1	Decode → Processor Control	Init state machine for instruction type and addressing mode
2-6	Processor Control → Registers	Populate scratch registers based on addressing mode.
Last Cycle	Processor Control → ALU	Execute
Last Cycle	Processor Control → Registers	Instruction Fetch

State Machines

Each {instruction_type, addressing_mode} triggers its own state machine. In brief, this state machine is responsible for signalling the Registers module to load/store addresses from memory or from the ALU.

State machine spec for each instruction type and addressing mode can be found at

https://docs.google.com/spreadsheets/d/16uGTSJEzrANUzr7dMmRNFAwA-_sEox-QsTjSI06IE/edit?usp=sharing

Considering one of the simplest instructions ADC immediate, which takes two cycles, the state machine is as follows:

Instruction_type=ARITHMETIC, addressing mode= IMMEDIATE

state=0	state=1	state=2
ld_sel=LD_INSTR; //instr= memory_data pc_sel=INC_PC; //pc++ next_state=state+1'b1	ld_sel=LD_IMM; //imm=memory_data pc_sel=INC_PC next_state=state+1'b1	alu_ctrl=DO_OP_ADC // execute src1_sel=SRC1_A src2_sel=SRC2_IMM dest_sel=DEST_A ld_sel=LD_INSTR//fetch next instruction pc_sel=INC_PC

		next_state=1'b1
--	--	-----------------

All instructions are classified into one of 55 state machines in the cpu specification sheet. The 6502 can take variable time for a single instructions based on certain conditions(page_cross, branch_taken etc). These corner case state transitions are also taken care of by processor control.

CPU Top Level Interface

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst	input		System active high reset
nmi	input	PPU	Non maskable interrupt from PPU. Executes BRK instruction in CPU
addr[15:0]	output	RAM	Address for R/W issued by CPU
dout[7:0]	input/output	RAM	Data from the RAM in case of reads and to the RAM in case of writes
memory_read	output	RAM	read enable signal for RAM
memory_write	output	RAM	write enable signal for RAM

CPU Instruction Decode Interface

The decode module is responsible for classifying the instruction into one of the addressing modes and an instruction type. It also generates the signal that the ALU would eventually use if the instruction passed through the ALU.

Signal name	Signal Type	Source/Dest	Description
instruction_register	input	Registers	Opcode of the current instruction
nmi	input	cpu_top	Non maskable interrupt from PPU. Executes BRK instruction in CPU
instruction_type	output	Processor Control	Type of instruction. Belongs to enum ITYPE.
addressing_mode	output	Processor Control	Addressing mode of the opcode in instruction_register. Belongs to enum AMODE.
alu_sel	output	ALU	ALU operation expected to be performed by the opcode, eventually. Processor control chooses to use it at a cycle appropriate for the instruction. Belongs to enum DO_OP.

CPU MEM Interface

The MEM module is the interface between memory and CPU. It provides appropriate address and read/write signal for the memory. Controlled by the select signals

Signal name	Signal Type	Source/Dest	Description
addr_sel	input	Processor Control	Selects which input to use as address to memory. Enum of ADDR
int_sel	input	Processor Control	Selects which interrupt address to jump to. Enum of INT_TYPE
ld_sel,st_sel	input	Processor Control	Decides whether to read or write based on these signals
ad, ba, sp, irql, irqh, pc	input	Registers	Registers that are candidates of the address
addr	output	Memory	Address of the memory to read/write
read,write	output	Memory	Selects whether Memory should read or write

CPU ALU Interface

Performs arithmetic, logical operations and operations that involve status registers.

Signal name	Signal Type	Source/Dest	Description
in1, in2	input	ALU Input Selector	Inputs to the ALU operations selected by ALU Input module.
alu_sel	input	Processor Control	ALU operation expected to be performed by the opcode, eventually. Processor control chooses to use it at a cycle appropriate for the instruction. Belongs to enum DO_OP.
clk, rst	input		System clock and active high reset
out	output	to all registers	Output of ALU operation. sent to all registers and registers decide whether to receive it or ignore it as its next value.
n, z, v, c, b, d, i	output		Status Register

ALU Input Selector

Selects the input1 and input2 for the ALU

Signal name	Signal Type	Source/Dest	Description
src1_sel, src2_sel	input	Processor Control	Control signal that determines which sources to take in as inputs to ALU according to the instruction and addressing mode
a, bal, bah, adl, pcl, pch, imm, adv, x, bav, y, offset	input	Registers	Registers that are candidates to the input to ALU
temp_status	input	ALU	Sometimes status information is required but we don't want it to affect the status register. So we directly receive temp_status value from ALU
in1, in2	output	ALU	Selected input for the ALU

CPU Registers Interface

Holds all of the registers

Signal name	Signal Type	Source/Dest	Description
clk, rst	input		System clk and rst
dest_sel, pc_sel, sp_sel, ld_sel, st_sel	input	Processor Control	Selects which input to accept as new input. enum of DEST, PC, SP, LD, ST
clr_adh, clr_bah	input	Processor Control	Clears the high byte of ad, ba
alu_out, next_status	input	ALU	Output from ALU and next status value. alu_out can be written to most of the registers
data	inout	Memory	Datapath to Memory. Either receives or sends data according to ld_sel and st_sel.
a, x, y, ir, imm, adv, bav, offset, sp, pc, ad, ba, n, z, v,	output		Register outputs that can be used by different modules

c, b, d, i, status			
--------------------	--	--	--

CPU Processor Control Interface

The processor control module maintains the current state that the instruction is in and decides the control signals for the next state. Once the instruction type and addressing modes are decoded, the processor control block becomes aware of the number of cycles the instruction will take. Thereafter, at each clock cycle it generates the required control signals.

Signal name	Signal Type	Source/Dest	Description
instruction_type	input	Decode	Type of instruction. Belongs to enum ITYPE.
addressing_mode	input	Decode	Addressing mode of the opcode in instruction_register. Belongs to enum AMODE.
alu_ctrl	input	Decode	ALU operation expected to be performed by the opcode, eventually. Processor control chooses to use it at a cycle appropriate for the instruction. Belongs to enum DO_OP.
reset_adh	output	Registers	Resets ADH register
reset_bah	output	Registers	Resets BAH register
set_b	output	Registers	Sets the B flag
addr_sel	output	Registers	Selects the value that needs to be set on the address bus. Belongs to enum ADDR
alu_sel	output	ALU	Selects the operation to be performed by the ALU in the current cycle. Belongs to enum DO_OP
dest_sel	output	Registers	Selects the register that receives the value from ALU output. Belongs to enum DEST
ld_sel	output	Registers	Selects the register that will receive the value from Memory Bus. Belongs to enum LD
pc_sel	output	Registers	Selects the value that the PC will take next cycle. Belongs to enum PC
sp_sel	output	Registers	Selects the value that the SP will take next cycle. Belongs to enum SP

src1_sel	output	ALU	Selects src1 for ALU. Belongs to enum SRC1
src2_sel	output	ALU	Selects src2 for ALU. Belongs to enum SRC2
st_sel	output	Registers	Selects the register whose value will be placed on dout. Belongs to enum ST

CPU Enums

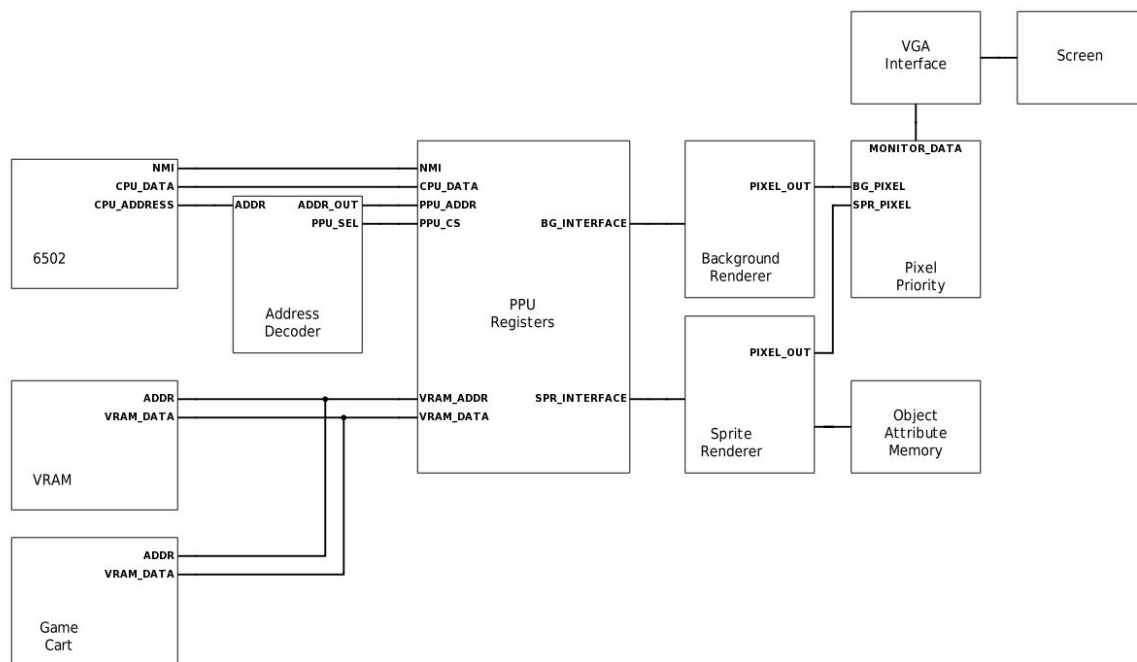
Enum name	Legal Values
ITYPE	ARITHMETIC,BRANCH,BREAK,CMPLDX,CMPLDY,INTERRUPT,JSR,JUMP,OTHER,PULL,PUSH,RMW,RTI,RTS,STA,STX,STY
AMODE	ABSOLUTE,ABSOLUTE_INDEX,ABSOLUTE_INDEX_Y,ACCUMULATOR,IMMEDIATE,IMPLIED,INDIRECT,INDIRECT_X,INDIRECT_Y,RELATIVE,SPECIAL,ZEROPAGE,ZEROPAGE_INDEX,ZEROPAGE_INDEX_Y
DO_OP	DO_OP_ADD,DO_OP_SUB,DO_OP_AND,DO_OP_OR,DO_OP_XOR,DO_OP_ASL,DO_OP_LSR,DO_OP_ROL,DO_OP_ROR,DO_OP_SRC2DO_OP_CLR_C,DO_OP_CLR_I,DO_OP_CLR_V,DO_OP_SET_C,DO_OP_SET_I,DO_OP_SET_V
ADDR	ADDR_AD,ADDR_PC,ADDR_BA,ADDR_SP,ADDR_IRQI,ADDR_IRQH
LD	LD_INSTR,LD_ADL,LD_ADH,LD_BAL,LD_BAH,LD_IMM,LD_OFFSET,LD_ADV,LD_BAV,LD_PCL,LD_PCH
SRC1	SRC1_A,SRC1_BAL,SRC1_BAH,SRC1_ADL,SRC1_PCL,SRC1_PCH,SRC1_BAV,SRC1_1
SRC2	SRC2_DC,SRC2_IMM,SRC2_ADV,SRC2_X,SRC2_BAV,SRC2_C,SRC2_1,SRC2_Y,SRC2_OFFSET
DEST	DEST_BAL,DEST_BAH,DEST_ADL,DEST_A,DEST_X,DEST_Y,DEST_PCL,DEST_PCH,DEST_NONE
PC	AD_P_TO_PC,INC_PC,KEEP_PC
SP	INC_SP,DEC_SP

4. Picture Processing Unit

The NES picture processing unit or PPU is the unit responsible for handling all of the console's graphical workloads. Obviously this is useful to the CPU because it offloads the highly intensive task of rendering a frame. This means the CPU can spend more time performing the game logic.

The PPU renders a frame by reading in scene data from various memories the PPU has access to such as VRAM, the game cart, and object attribute memory and then outputting an NTSC compliant 256x240 video signal at 60 Hz. The PPU was a special custom designed IC for Nintendo, so no other devices use this specific chip. It operates at a clock speed of 5.32 MHz making it three times faster than the NES CPU. This is one of the areas of difficulty in creating the PPU because it is easy to get the CPU and PPU clock domains out of sync.

PPU Top Level Schematic



6502: The CPU used in the NES. Communicates with the PPU through simple load/store instructions. This works because the PPU registers are memory mapped into the CPU's address space.

Address Decoder: The address decoder is responsible for selecting the chip select of the device the CPU wants to talk to. In the case of the PPU the address decoder will activate if from addresses [0x2000,

0x2007].

VRAM: The PPU video memory contains the data needed to render the scene, specifically it holds the name tables. VRAM is 2 Kb in size and depending on how the PPU VRAM address lines are configured, different mirroring schemes are possible.

Game Cart: The game cart has a special ROM on it called the character ROM, or char ROM for short. the char ROM contains the sprite and background tile pattern data. These are sometimes referred to as the pattern tables.

PPU Registers: These registers allow the CPU to modify the state of the PPU. It maintains all of the control signals that are sent to both the background and sprite renderers.

Background Renderer: Responsible for drawing the background data for a scene.

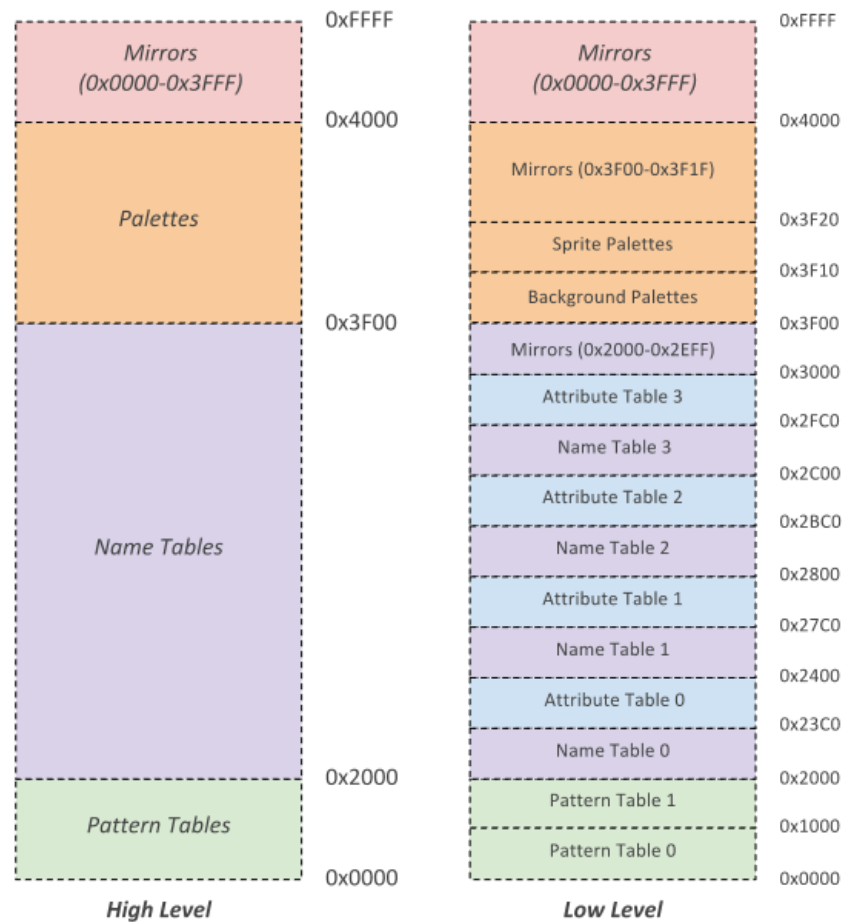
Sprite Renderer: Responsible for drawing the sprite data for a scene, and maintaining object attribute memory.

Object Attribute Memory: Holds all of the data needed to know how to render a sprite. OAM is 256 bytes in size and each sprite utilizes 4 bytes of data. This means the PPU can support 64 sprites.

Pixel Priority: During the visible pixel section of rendering, both the background and sprite renderers produce a single pixel each clock cycle. The pixel priority module looks at the priority values and color for each pixel and decides which one to draw to the screen.

VGA Interface: This is where all of the frame data is kept in a frame buffer. This data is then upscaled to 640x480 when it goes out to the monitor.

PPU Memory Map



The PPU memory map is broken up into three distinct regions, the pattern tables, name tables, and palettes. Each of these regions holds data the PPU need to obtain to render a given scanline. The functionality of each part is described in the PPU Rendering section.


PPU CHAROM

- ROM from the cartridge is broken in two sections
 - Program ROM
 - Contains program code for the 6502
 - Is mapped into the CPU address space by the mapper
 - Character ROM
 - Contains sprite and background data for the PPU
 - Is mapped into the PPU address space by the mapper


PPU Rendering

- Pattern Tables
 - \$0000-\$2000 in VRAM
 - Pattern Table 0 (\$0000-\$0FFF)
 - Pattern Table 1 (\$1000-\$2000)
 - The program selects which one of these contains sprites and backgrounds
 - Each pattern table is 16 bytes long and represents 1 8x8 pixel tile

- Each 8x1 row is 2 bytes long
- Each bit in the byte represents a pixel and the corresponding bit for each byte is combined to create a 2 bit color.
 - Color_pixel = {byte2[0], byte1[0]}
- So there can only be 4 colors in any given tile
- Rightmost bit is leftmost pixel
- Any pattern that has a value of 0 is transparent i.e. the background color



- Name Tables
 - \$2000-\$2FFF in VRAM with \$3000-\$3EFF as a mirror
 - Laid out in memory in 32x30 fashion
 - Think of as a 2d array where each element specifies a tile in the pattern table.
 - This results in a resolution of 256x240
 - Although the PPU supports 4 name tables the NES only supplied enough VRAM for 2 this results in 2 of the 4 name tables being mirror
 - Vertically = horizontal movement
 - Horizontally = vertical movement
 - Each entry in the name table refers to one pattern table and is one byte. Since there are 32x30=960 entries each name table requires 960 bytes of space the left over 64 bytes are used for attribute tables
 - Attribute tables
 - 1 byte entries that contains the palette assignment for a 2x2 grid of tiles



- Sprites
 - Just like backgrounds sprite tile data is contained in one of the pattern tables
 - But unlike backgrounds sprite information is not contained in name tables but in a special reserved 256 byte RAM called the object attribute memory (OAM)
- Object Attribute Memory
 - 256 bytes of dedicated RAM
 - Each object is allocated 4 bytes of OAM so we can store data about 64 sprites at once
 - Each object has the following information stored in OAM
 - X Coordinate
 - Y Coordinate
 - Pattern Table Index
 - Palette Assignment
 - Horizontal/Vertical Flip
- Palette Table
 - Located at \$3F00-\$3F20
 - \$3F00-\$3F0F is background palettes
 - \$3F10-\$3F1F is sprite palettes
 - Mirrored all the way to \$4000
 - Each color takes one byte
 - Every background tile and sprite needs a color palette.
 - When the background or sprite is being rendered the the color for a specific table is looked up in the correct palette and sent to the draw select mux.
- Rendering is broken into two parts which are done for each horizontal scanline
 - Background Rendering
 - The background enable register (\$2001) controls if the default background color is rendered (\$2001) or if background data from the background renderer.
 - The background data is obtained for every pixel.
 - Sprite Rendering
 - The sprite renderer has room for 8 unique sprites on each scanline.
 - For each scanline the renderer looks through the OAM for sprites that need to be drawn on the scanline. If this is the case the sprite is loaded into the scanline local sprites
 - If this number exceeds 8 a flag is set and the behavior is undefined.
 - If a sprite should be drawn for a pixel instead of the background the sprite renderer sets the sprite priority line to a mux that decides what to send to the screen and the mux selects the sprite color data.

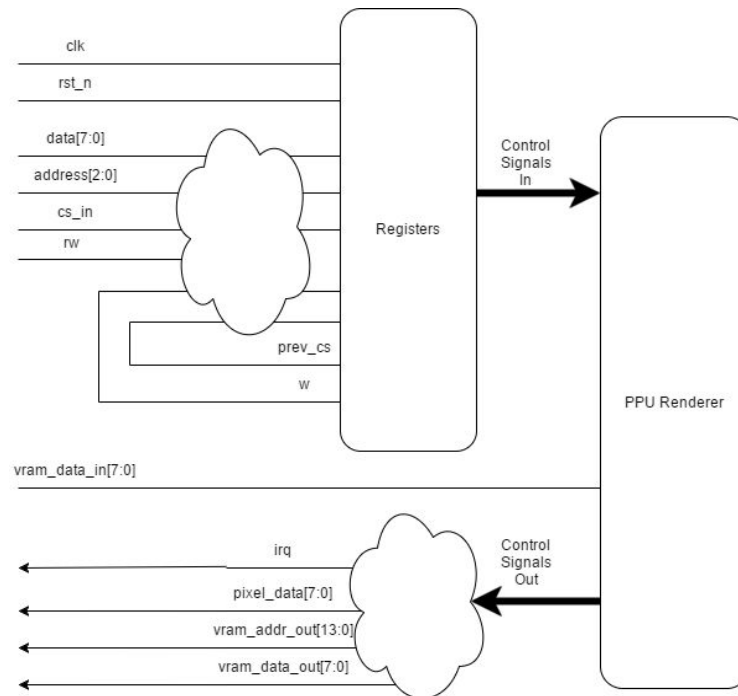
PPU Memory Mapped Registers

The PPU register interface exists so the CPU can modify and fetch the state elements of the PPU. These state elements include registers that set control signals, VRAM, object attribute memory, and palettes. These state elements then determine how the background and sprite renderers will draw the scene. The PPU register module also contains the pixel mux and palette memory which are used to determine what pixel data to send to the VGA module.

Signal name	Signal	Source/De	Description
-------------	--------	-----------	-------------

	Type	st	
clk	input		System clock
rst_n	input		System active low reset
data[7:0]	inout	CPU	Bi directional data bus between the CPU/PPU
address[2:0]	input	CPU	Register select
rw	input	CPU	CPU read/write select
cs_in	input	CPU	PPU chip select
irq	output	CPU	Signal PPU asserts to trigger CPU NMI
pixel_data[7:0]	output	VGA	RGB pixel data to be sent to the display
vram_addr_out[13:0]	output	VRAM	VRAM address to read/write from
vram_rw_sel	output	VRAM	Determines if the current vram operation is a read or write
vram_data_out[7:0]	output	VRAM	Data to write to VRAM
frame_end	output	VGA	Signals the VGA interface that this is the end of a frame
frame_start	output	VGA	Signals the VGA interface that a frame is starting to keep the PPU and VGA in sync
rendering	output	VGA	Signals the VGA interface that pixel data output is valid

PPU Register Block Diagram



PPU Register Descriptions

- Control registers are mapped into the CPU's address space (\$2000 - \$2007)
- The registers are repeated every eight bytes until address \$3FF
- **PPUCTRL[7:0] (\$2000) WRITE**
 - [1:0]: Base nametable address which is loaded at the start of a frame
 - 0: \$2000
 - 1: \$2400
 - 2: \$2800
 - 3: \$2C00
 - [2]: VRAM address increment per CPU read/write of PPUDATA
 - 0: Add 1 going across
 - 1: Add 32 going down
 - [3]: Sprite pattern table for 8x8 sprites
 - 0: \$0000
 - 1: \$1000
 - Ignored in 8x16 sprite mode
 - [4]: Background pattern table address
 - 0: \$0000
 - 1: \$1000
 - [5]: Sprite size
 - 0: 8x8
 - 1: 8x16
 - [6]: PPU master/slave select
 - 0: Read backdrop from EXT pins
 - 1: Output color on EXT pins
 - [7]: Generate NMI interrupt at the start of vertical blanking interval
 - 0: off

- 1: on
- **PPUMASK[7:0]** (\$2001) WRITE
 - [0]: Use grayscale image
 - 0: Normal color
 - 1: Grayscale
 - [1]: Show left 8 pixels of background
 - 0: Hide
 - 1: Show background in leftmost 8 pixels of screen
 - [2]: Show left 8 pixels of sprites
 - 0: Hide
 - 1: Show sprites in leftmost 8 pixels of screen
 - [3]: Render the background
 - 0: Don't show background
 - 1: Show background
 - [4]: Render the sprites
 - 0: Don't show sprites
 - 1: Show sprites
 - [5]: Emphasize red
 - [6]: Emphasize green
 - [7]: Emphasize blue
- **PPUSTATUS[7:0]** (\$2002) READ
 - [4:0]: Nothing?
 - [5]: Set for sprite overflow which is when more than 8 sprites exist in one scanline (Is actually more complicated than this to do a hardware bug)
 - [6]: Sprite 0 hit. This bit gets set when a non zero part of sprite zero overlaps a non zero background pixel
 - [7]: Vertical blank status register
 - 0: Not in vertical blank
 - 1: Currently in vertical blank
- **OAMADDR[7:0]** (\$2003) WRITE
 - Address of the object attribute memory the program wants to access
- **OAMDATA[7:0]** (\$2004) READ/WRITE
 - The CPU can read/write this register to read or write to the PPU's object attribute memory. The address should be specified by writing the OAMADDR register beforehand. Each write will increment the address by one, but a read will not modify the address
- **PPUSCROLL[7:0]** (\$2005) WRITE
 - Tells the PPU what pixel of the nametable selected in PPUCTRL should be in the top left hand corner of the screen
- **PPUADDR[7:0]** (\$2006) WRITE
 - Address the CPU wants to write to VRAM before writing a read of PPUSTATUS is required and then two bytes are written in first the high byte then the low byte
- **PPUDATA[7:0]** (\$2007) READ/WRITE
 - Writes/Reads data from VRAM for the CPU. The value in PPUADDR is then incremented by the value specified in PPUCTRL
- **OAMDMA[7:0]** (\$4014) WRITE
 - A write of \$XX to this register will result in the CPU memory page at \$XX00-\$XXFF being written into the PPU object attribute memory

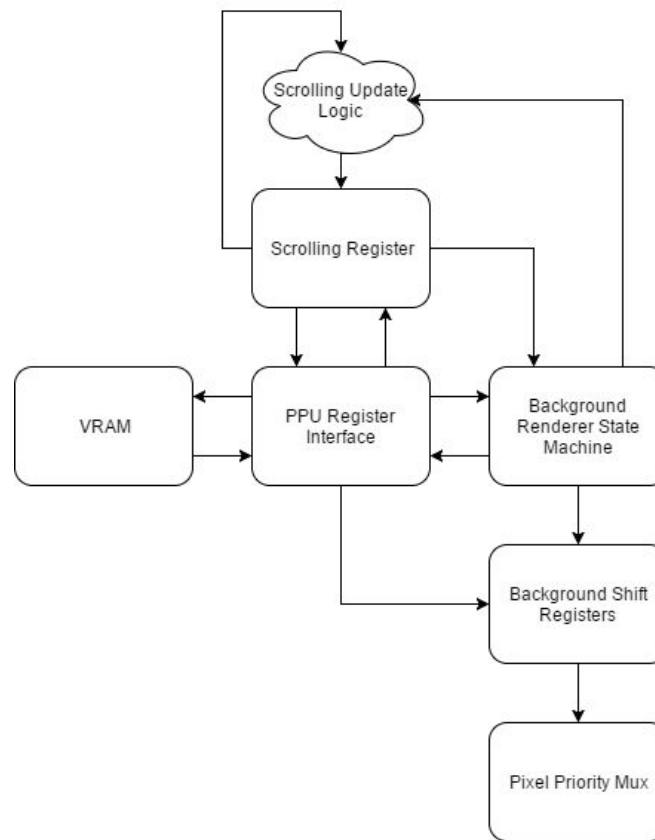
PPU Background Renderer

The background renderer is responsible for rendering the background for each frame that is output to the VGA interface. It does this by prefetching the data for two tiles at the end of the previous scanline. And then begins to continuously fetch tile data for every pixel of the visible frame. This allows the background renderer to produce a steady flow of output pixels despite the fact it takes 8 cycles to fetch 8 pixels of a scanline.

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
bg_render_en	input	PPU Register	Background render enable
x_pos[9:0]	input	PPU Register	The current pixel for the active scanline
y_pos[9:0]	input	PPU Register	The current scanline being rendered
vram_data_in[7:0]	input	PPU Register	The current data that has been read in from VRAM
bg_pt_sel	input	PPU Register	Selects the location of the background renderer pattern table
show_bg_left_col	input	PPU Register	Determines if the background for the leftmost 8 pixels of each scanline will be drawn
fine_x_scroll[2:0]	input	PPU Register	Selects the pixel drawn on the left hand side of the screen
coarse_x_scroll[4:0]	input	PPU Register	Selects the tile to start rendering from in the x direction
fine_y_scroll[2:0]	input	PPU Register	Selects the pixel drawn on the top of the screen
coarse_y_scroll[4:0]	input	PPU Register	Selects the tile to start rendering from in the y direction
nametable_sel[1:0]	input	PPU Register	Selects the nametable to start rendering from
update_loopy_v	input	PPU Register	Signal to update the temporary vram address
cpu_loopy_v_inc	input	PPU Register	Signal to increment the temporary vram address by the increment amount
cpu_loopy_v_inc_amt	input	PPU Register	If this signal is set increment the temp vram address by 32 on cpu_loopy_v_inc, and increment by 1 if it is not set on cpu_loopy_v_inc
vblank_out	output	PPU Register	Determines if the PPU is in vertical blank

bg_rendering_out	output	PPU Register	Determines if the bg renderer is requesting vram reads
bg_pal_sel[3:0]	output	Pixel Mux	Selects the palette for the background pixel
loopy_v_out[14:0]	output	PPU Register	The temporary vram address register for vram reads/writes
vram_addr_out[13:0]	output	VRAM	The VRAM address the sprite renderer wants to read from

PPU Background Renderer Diagram



VRAM: The background renderer reads from two of the three major areas of address space available to the PPU, the pattern tables, and the name tables. First the background renderer needs the name table byte for a given tile to know which tile to draw in the background. Once it has this information it needs the pattern to know how to draw the background tile.

PPU Register Interface: All background rendering VRAM reads are performed through the PPU register interface. This allows for vram address bus arbitration between the background renderer, sprite renderer, and the cpu.

Scrolling Register: The scrolling register is responsible for keeping track of what tile is currently being drawn to the screen.

Scrolling Update Logic: Every time the data for a background tile is successfully fetched the scrolling register needs to be updated. Most of the time it is a simple increment, but more care has to be taken when the next tile falls in another name table. This logic allows the scrolling register to correctly update to be able to smoothly jump between name tables while rendering.

Background Renderer State Machine: The background renderer state machine is responsible for sending the correct control signals to all of the other modules as the background is rendering.

Background Shift Registers: These registers shift out the pixel data to be rendered on every clock cycle. They also implement the logic that makes fine one pixel scrolling possible by changing what index of the shift registers is the one being shifted out each cycle.

Pixel Priority Mux: Since both the sprite renderer and background renderer output one pixel every clock cycle during the visible part of the frame, there needs to be some logic to pick between the two pixels that are output. The pixel priority mux does this based on the priority of the sprite pixel, and the color of both the sprite pixel and background pixel.

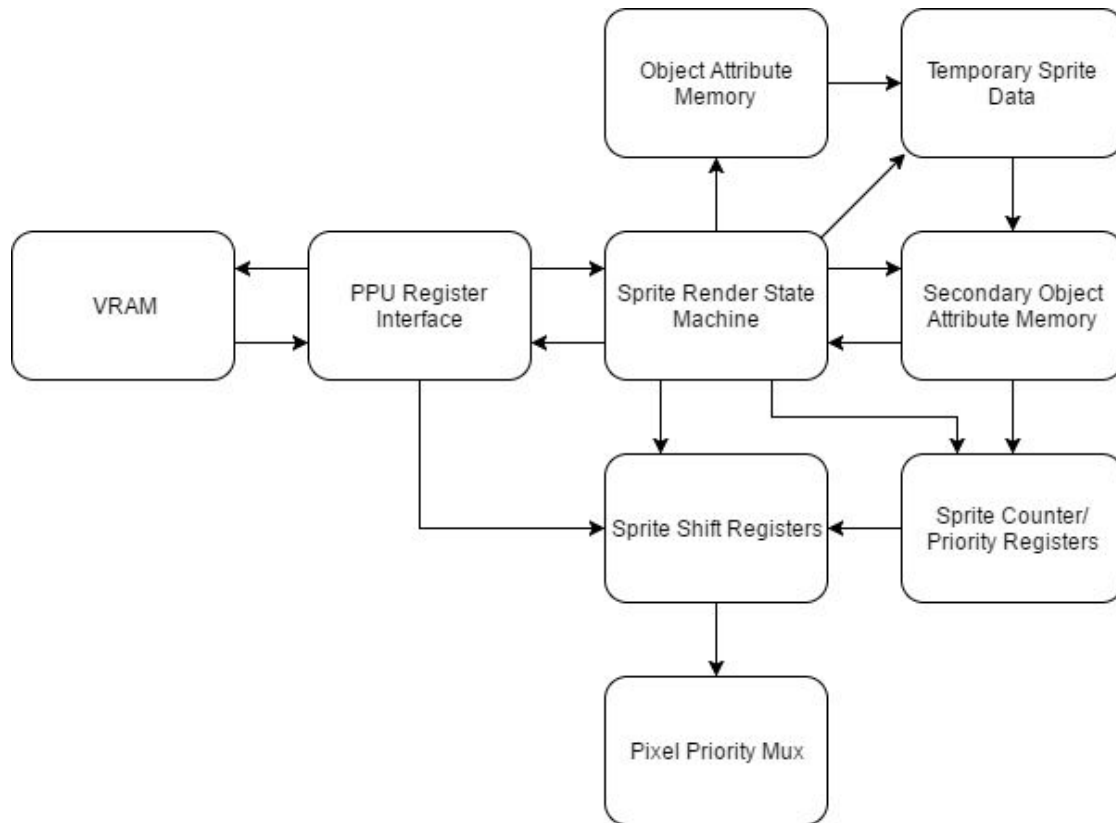
PPU Sprite Renderer

The PPU sprite renderer is used to render all of the sprite data for each scanline. The way the hardware was designed it only allows for 64 sprites to be kept in object attribute memory at once. There are only 8 spots available to store the sprite data for each scanline so only 8 sprites can be rendered for each scanline. Sprite data in OAM is evaluated for the next scanline while the background renderer is mastering the VRAM bus. When rendering reaches horizontal blank the sprite renderer fetches the pattern data for all of the sprites to be rendered on the next scanline and places the data in the sprite shift registers. The sprite x position is also loaded into a down counter which determines when to make the sprite active and shift out the pattern data on the next scanline.

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
spr_render_en	input	PPU Register	Sprite renderer enable signal
x_pos[9:0]	input	PPU Register	The current pixel for the active scanline
y_pos[9:0]	input	PPU Register	The current scanline being rendered
spr_addr_in[7:0]	input	PPU Register	The current OAM address being read/written
spr_data_in[7:0]	inout	PPU Register	The current data being read/written from OAM
vram_data_in[7:0]	input	VRAM	The data the sprite renderer requested from VRAM
cpu_oam_rw	input	PPU Register	Selects if OAM is being read from or written to from the CPU

cpu_oam_req	input	PPU Register	Signals the CPU wants to read/write OAM
spr_pt_sel	input	PPU Register	Determines the PPU pattern table address in VRAM
spr_size_sel	input	PPU Register	Determines the size of the sprites to be drawn
show_spr_left_col	input	PPU Register	Determines if sprites on the leftmost 8 pixels of each scanline will be drawn
spr_overflow	output	PPU Register	If more than 8 sprites fall on a single scanline this is set
spr_pri_out	output	Pixel Mux	Determines the priority of the sprite pixel data
spr_data_out[7:0]	output	PPU Register	returns oam data the CPU requested
spr_pal_sel[3:0]	output	Pixel Mux	Sprite pixel color data to be drawn
vram_addr_out[13:0]	output	VRAM	Sprite vram read address
spr_vram_req	output	VRAM	Signals the sprite renderer is requesting a VRAM read
spr_0_rendering	output	Pixel Mux	Determines if the current sprite that is rendering is sprite 0
inc_oam_addr	output	PPU Register	Signals the OAM address in the registers to increment

PPU Sprite Renderer Diagram



VRAM: The sprite renderer needs to be able to fetch the sprite pattern data from the character rom. This is why it can request VRAM reads from this region through the PPU Register Interface

PPU Register Interface: All background rendering VRAM reads are performed through the PPU register interface. This allows for vram address bus arbitration between the background renderer, sprite renderer, and the cpu.

Object Attribute Memory: OAM contains all of the data needed to render a sprite to the screen except the pattern data itself. For each sprite OAM holds its x position, y position, horizontal flip, vertical flip, and palette information. In total OAM supports 64 sprites.

Sprite Renderer State Machine: The sprite renderer state machine is responsible for sending all of the control signals to each of the other units in the renderer. This includes processing the data in OAM, moving the correct sprites to secondary OAM, VRAM reads, and shifting out the sprite data when the sprite needs to be rendered to the screen.

Sprite Shift Registers: The sprite shift registers hold the sprite pixel data for sprites on the current scanline. When a sprite becomes active its data is shifted out to the pixel priority mux.

Pixel Priority Mux: Since both the sprite renderer and background renderer output one pixel every clock cycle during the visible part of the frame, there needs to be some logic to pick between the two pixels that are output. The pixel priority mux does this based on the priority of the sprite pixel, and the color of both the sprite pixel and background pixel.

Temporary Sprite Data: The temporary sprite data is where the state machine moves the current sprite being evaluated in OAM to. If the temporary sprite falls on the next scanline its data is moved into a slot in secondary OAM. If it does not the data is discarded.

Secondary Object Attribute Memory: Secondary OAM holds the sprite data for sprites that fall on the next scanline. During hblank this data is used to load the sprite shift registers with the correct sprite pattern data.

Sprite Counter and Priority Registers: These registers hold the priority information for each sprite in the sprite shift registers. It also holds a down counter for each sprite which is loaded with the sprite's x position. When the counter hits 0 the corresponding sprite becomes active and the sprite data needs to be shifted out to the screen.

PPU Object Attribute Memory

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
oam_en	input	OAM	Determines if the input data is for a valid read/write
oam_rw	input	OAM	Determines if the current operation is a read or write
spr_select[5:0]	input	OAM	Determines which sprite is being read/written
byte_select[1:0]	input	OAM	Determines which sprite byte is being read/written
data_in[7:0]	input	OAM	Data to write to the specified OAM address
data_out[7:0]	output	PPU Register	Data that has been read from OAM

PPU Palette Memory

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
pal_addr[4:0]	input	palette mem	Selects the palette to read/write in the memory
pal_data_in[7:0]	input	palette mem	Data to write to the palette memory
palette_mem_rw	input	palette mem	Determines if the current operation is a read or write
palette_mem_en	input	palette mem	Determines if the palette mem inputs are valid
color_out[7:0]	output	VGA	Returns the selected palette for a given address on a read

VRAM Interface

The VRAM interface instantiates an Altera RAM IP core. Each read take 2 cycles one for the input and one for the output

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
vram_addr[10:0]	input	PPU	Address from VRAM to read to or write from
vram_data_in[7:0]	input	PPU	The data to write to VRAM
vram_en	input	PPU	The VRAM enable signal
vram_rw	input	PPU	Selects if the current op is a read or write
vram_data_out[7:0]	output	PPU	The data that was read from VRAM on a read

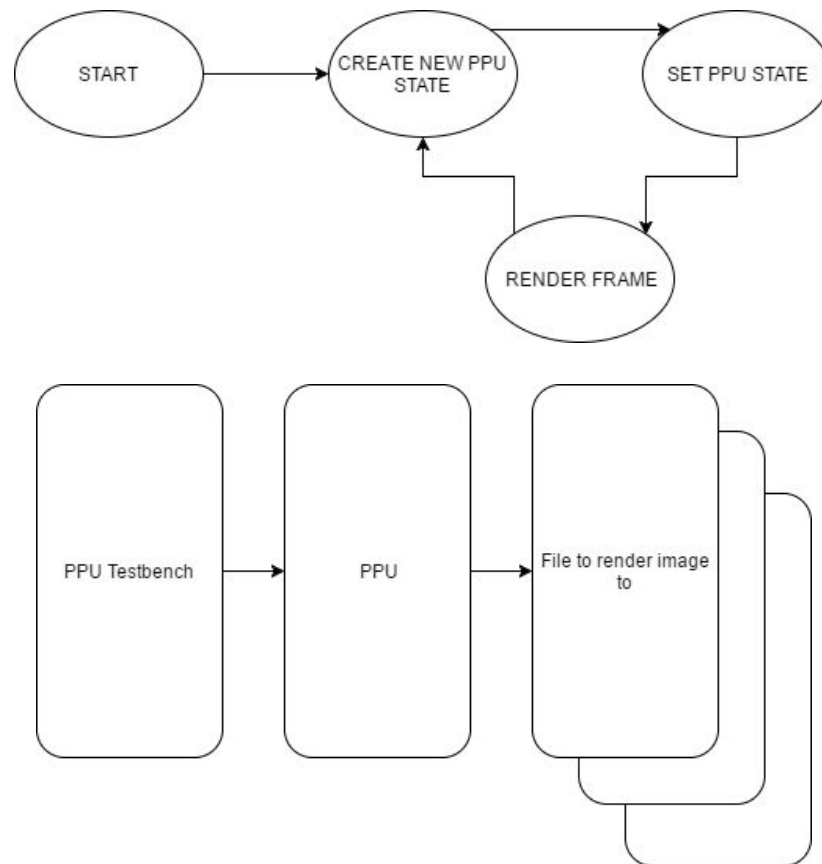
DMA

The DMA is used to copy 256 bytes of data from the CPU address space into the OAM (PPU address space). The DMA is 4x faster than it would be to use str and ldr instructions to copy the data. While copying data, the CPU is stalled.

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
oamdma	input	PPU	When written to, the DMA will begin copying data to the OAM. If the value written here is XX then the data that will be copied begins at the address XX00 in the CPU RAM and goes until the address XXFF. Data will be copied to the OAM starting at the OAM address specified in the OAMADDR register of the OAM.
cpu_ram_q	input	CPU RAM	Data read in from CPU RAM will come here
dma_done	output	CPU	Informs the CPU to pause while the DMA copies OAM data from the CPU RAM to the OAM section of the PPU RAM
cpu_ram_addr	output	CPU RAM	The address of the CPU RAM where we are reading data
cpu_ram_wr	output	CPU RAM	Read/write enable signal for CPU RAM
oam_data	output	OAM	The data that will be written to the OAM at the address specified in OAMADDR
dma_req	input	APU	High when the DMC wants to use the DMA
dma_ack	output	APU	High when data on DMA
dma_addr	input	APU	Address for DMA to read from ** CURRENTLY NOT USED **
dma_data	output	APU	Data from DMA to apu memory ** CURRENTLY NOT USED **

PPU Testbench

In a single frame the PPU outputs 61,440 pixels. Obviously this amount of information would be incredibly difficult for a human to verify as correct by looking at a simulation waveform. This is what drove me to create a testbench capable of rendering full PPU frames to an image. This allowed the process of debugging the PPU to proceed at a much faster rate than if I used waveforms alone. Essentially how the test bench works is the testbench sets the initial PPU state, it lets the PPU render a frame, and then the testbench receives the data for each pixel and generates a PPM file. The testbench can render multiple frames in a row, so the tester can see how the frame output changes as the PPU state changes.



PPU Testbench PPM file format

The PPM image format is one of the easiest to understand image formats available. This is mostly because of how it is a completely human readable format. A PPM file simply consists of a header, and then pixel data. The header consists of the text "P3" on the first line, followed by the image width and height on the next line, then a max value for each of the rgb components of a pixel on the final line of the header. After the header it is just width * height rgb colors in row major order.

PPU Testbench Example Renderings



5. Memory Maps

Cartridges are a Read-Only Memory that contains necessary data to run games. However, it is possible that in some cases that a cartridge holds more data than the CPU can address to. In this case, memory mapper comes into play and changes the mapping as needed so that one address can point to multiple locations in a cartridge. For our case, the end goal was to get the game Super Mario Bros. running on our FPGA. This game does not use a memory mapper, so we did not work on any memory mappers. In the future, we might add support for the other memory mapping systems so that we can play other games.

These were two ip catalog ROM blocks that are created using MIF files for Super Mario Bros. They contained the information for the CPU and PPU RAM and VRAM respectively.

PPU ROM Memory Map

This table shows how the PPU's memory is laid out. The Registers are explained in greater detail in the Architecture Document.

Address Range	Description
0x0000 - 0x0FFF	Pattern Table 0
0x1000 - 0x1FFF	Pattern Table 1
0x2000 - 0x23BF	Name Table 0
0x23C0 - 0x23FF	Attribute Table 0
0x2400 - 0x27BF	Name Table 1
0x27C0 - 0x27FF	Attribute Table 1
0x2800 - 0x2BBF	Name Table 2
0x2BC0 - 0x2BFF	Attribute Table 2
0x2C00 - 0x2FBF	Name Table 3
0x2FC0 - 0x2FFF	Attribute Table 3
0x3000 - 0x3EFF	Mirrors 0x2000 - 0x2EFF
0x3F00 - 0x3F0F	Background Palettes
0x3F10 - 0x3F1F	Sprite Palettes
0x3F20 - 0x3FFF	Mirrors 0x3F00 - 0x3F1F
0x4000 - 0xFFFF	Mirrors 0x0000 - 0x3FFF

CPU ROM Memory Map

This table explains how the CPU's memory is laid out. The Registers are explained in greater detail in the Architecture document.

Address Range	Description
0x0000 - 0x00FF	Zero Page
0x0100 - 0x1FF	Stack
0x0200 - 0x07FF	RAM
0x0800 - 0x1FFF	Mirrors 0x0000 - 0x07FF
0x2000 - 0x2007	Registers

0x2008 - 0x3FFF	Mirrors 0x2000 - 0x2007
0x4000 - 0x401F	I/O Registers
0x4020 - 0x5FFF	Expansion ROM
0x6000 - 0x7FFF	SRAM
0x8000 - 0xBFFF	Program ROM Lower Bank
0xC000 - 0xFFFF	Program ROM Upper Bank

Memory Mappers Interface

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
rd	input	CPU/PPU	Read request
addr	input	CPU/PPU	Address to read from
data	output	CPU/PPU	Data from the address

6. APU

Due to limitations of our FPGA design board (no D2A converter) and time constraints, our group did not implement the APU. Instead, we created the register interface for the APU, so that the CPU could still read and write from the registers. The following section is provided for reference only.

The NES included an Audio Processing Unit (APU) to control all sound output. The APU contains five audio channels: two pulse wave modulation channels, a triangle wave channel, a noise channel (for random audio), and a delta modulation channel. Each channel is mapped to registers in the CPU's address space and each channel runs independently of the others. The outputs of all five channels are then combined using a non-linear mixing scheme. The APU also has a dedicated APU Status register. A write to this register can enable/disable any of the five channels. A read to this register can tell you if each channel still has a positive count on their respective timers. In addition, a read to this register will reveal any DMC or frame interrupts.

APU Registers

Registers			
\$4000	First pulse wave	DDLC VVVV	Duty, Envelope Loop, Constant Volume, Volume
\$4001	First pulse wave	EPPP NSSS	Enabled, Period, Negate, Shift
\$4002	First pulse wave	TTTT TTTT	Timer low
\$4003	First pulse wave	LLLL LTTT	Length counter load, Timer high
\$4004	Second pulse wave	DDLC VVVV	Duty, Envelope Loop, Constant Volume, Volume
\$4005	Second pulse wave	EPPP NSSS	Enabled, Period, Negate, Shift
\$4006	Second pulse wave	TTTT TTTT	Timer low
\$4007	Second pulse wave	LLLL LTTT	Length counter load, Timer high
\$4008	Triangle wave	CRRR RRRR	Length counter control, linear count load
\$4009	Triangle wave		Unused
\$400A	Triangle wave	TTTT TTTT	Timer low
\$400B	Triangle wave	LLLL LTTT	Length counter load, Timer high
\$400C	Noise Channel	--LC VVVV	Envelope Loop, Constant Volume, Volume
\$400D	Noise Channel		Unused
\$400E	Noise Channel	L--- PPPP	Loop Noise, Noise Period
\$400F	Noise Channel	LLLL L---	Length counter load
\$4010	Delta modulation channel	IL-- FFFF	IRQ enable, Loop, Frequency

\$4011	Delta modulation channel	-LLL LLLL	Load counter
\$4012	Delta modulation channel	AAAA AAAA	Sample Address
\$4013	Delta modulation channel	LLLL LLLL	Sample Length
\$4015 (write)	APU Status Register Writes	---D NT21	Enable DMC, Enable Noise, Enable Triangle, Enable Pulse 2/1
\$4015 (read)	APU Status Register Read	IF-D NT21	DMC Interrupt, Frame Interrupt, DMC Active, Length Counter > 0 for Noise, Triangle, and Pulse Channels
\$4017 (write)	APU Frame Counter	MI-- ----	Mode (0 = 4 step, 1 = 5 step), IRQ inhibit flag

7. Controllers (SPART)

The controller module allows users to provide input to the FPGA. We opted to create a controller simulator program instead of using an actual NES joystick. This decision was made because the NES controllers used a proprietary port and because the available USB controllers lacked specification sheets. The simulator program communicates with the FPGA using the SPART interface, which is similar to UART. Our SPART module used 8 data bits, no parity, and 1 stop bit for serial communication. All data was received automatically into an 8 bit buffer by the SPART module at 2400 baud. In addition to the SPART module, we also needed a controller driver to allow the CPU to interface with the controllers. The controllers are memory mapped to \$4016 and \$4017 for CPU to read.

When writing high to address \$4016 bit 0, the controllers are continuously loaded with the states of each button. Once address \$4016 bit 0 is cleared, the data from the controllers can be read by reading from address \$4016 for player 1 or \$4017 for player 2. The data will be read in serially on bit 0. The first read will return the state of button A, then B, Select, Start, Up, Down, Left, Right. It will read 1 if the button is pressed and 0 otherwise. Any read after clearing \$4016 bit 0 and after reading the first 8 button values, will be a 1. If the CPU reads when before clearing \$4016, the state of button A will be repeatedly returned.

Debug Modification

In order to provide an easy way to debug our top level design, we modified the controller to send an entire ram block out over SPART when it receives the send_states signal. This later allowed us to record the PC, IR, A, X, Y, flags, and SP of the CPU into a RAM block every CPU clock cycle and print this out onto a terminal console when we reached a specific PC.


Controller Registers

Registers			
\$4016 (write)	Controller Update	---- --C	Button states of both controllers are loaded
\$4016 (read)	Controller 1 Read	---- --C	Reads button states of controller 1 in the order A, B, Start, Select, Up, Down, Left, Right
\$4017 (read)	Controller 2 Read	---- --C	Reads button states of controller 2 in the order A, B, Start, Select, Up, Down, Left, Right

Controllers Wrapper

The controllers wrapper acts as the top level interface for the controllers. It instantiates two Controller modules and connects each one to separate TxD RxD lines. In addition, the Controllers wrapper handles passing the cs, addr, and rw lines into the controllers correctly. Both controllers receive an address of 0 for controller writes, while controller 1 will receive address 0 for reads and controller 2 will receive address 1.

Controller Wrapper Diagram



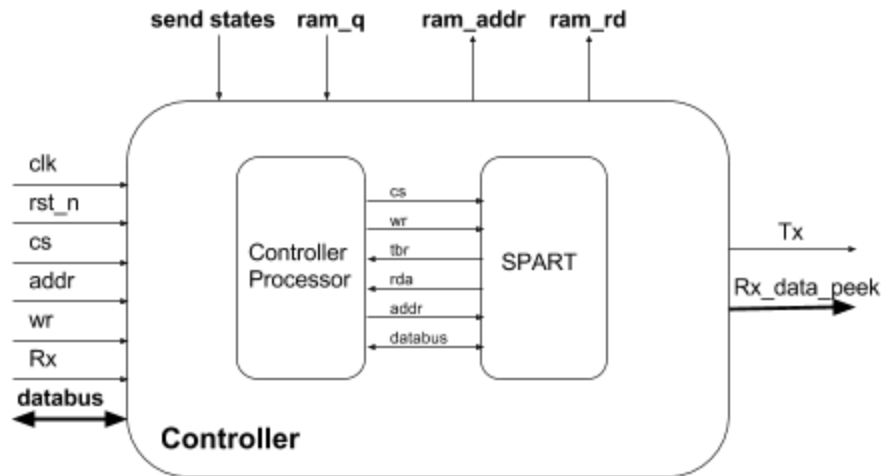
Controller Wrapper Interface

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
TxD1	output	UART	Transmit data line for controller 1
TxD2	output	UART	Transmit data line for controller 2
RxD1	input	UART	Receive data line for controller 1
RxD2	input	UART	Receive data line for controller 2
addr	input	CPU	Controller address, 0 for \$4016, 1 for \$4017
cpubus[7:0]	inout	CPU	Data from/to the CPU
cs	input	CPU	Chip select
rw	input	CPU	Read/Write signal (high for reads)
rx_data_peek	output	LEDR[7:0]	Output states to the FPGA LEDs to show that input was being received
send_states	input		When this signal goes high, the controller begins outputting RAM data
cpuram_q	input	CPU RAM	Stored CPU states from the RAM block
rd_addr	output	CPU RAM	The address the controller is writing out to SPART
rd	output	CPU RAM	High when controller is reading from CPU RAM

Controller

The controller module instantiates the Driver and SPART module's.

Controller Diagram



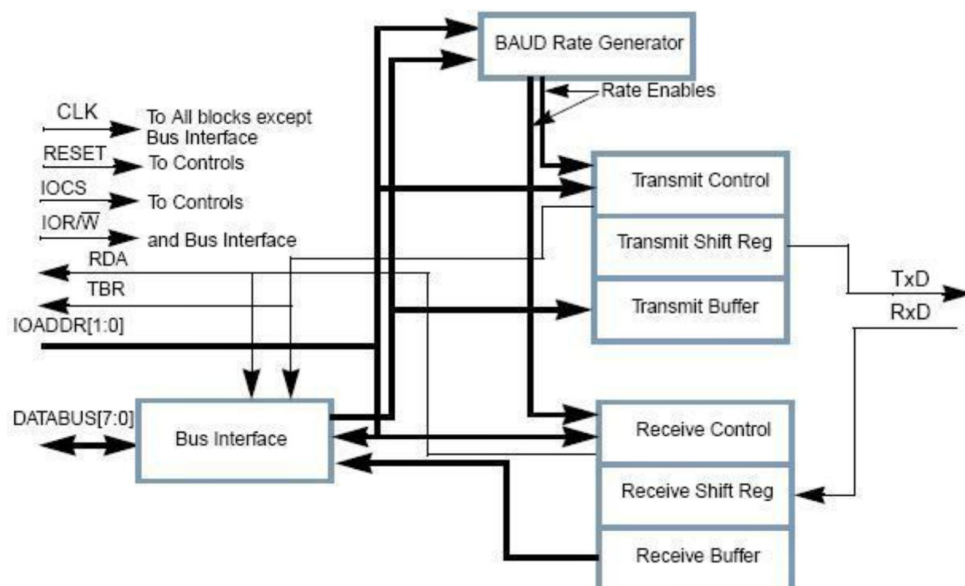
Controller Interface

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
TxD	output	UART	Transmit data line
RxD	input	UART	Receive data line
addr	input	CPU	Controller address 0 for \$4016, 1 for \$4017
dout[7:0]	inout	CPU	Data from/to the CPU
cs	input	CPU	Chip select
rw	input	CPU	Read write signal (low for writes)
rx_data_peek	output	LEDR	Outputs button states to FPGA LEDs
send_states	input		When this signal goes high, the controller begins outputting RAM data
cpuram_q	input	CPU RAM	Stored CPU states from the RAM block
rd_addr	output	CPU RAM	The address the controller is writing out to SPART
rd	output	CPU RAM	High when controller is reading from CPU RAM

Special Purpose Asynchronous Receiver and Transmitter (SPART)

The SPART Module is used to receive serial data. The SPART and driver share many interconnections in order to control the reception and transmission of data. On the left, the SPART interfaces to an 8-bit, 3-state bidirectional bus, DATABUS[7:0]. This bus is used to transfer data and control information between the driver and the SPART. In addition, there is a 2-bit address bus, IOADDR[1:0] which is used to select the particular register that interacts with the DATABUS during an I/O operation. The IOR/W signal determines the direction of data transfer between the driver and SPART. For a Read (IOR/W=1), data is transferred from the SPART to the driver and for a Write (IOR/W=0), data is transferred from the driver to the SPART. IOCS and IOR/W are crucial signals in properly controlling the three-state buffer on DATABUS within the SPART. Receive Data Available (RDA), is a status signal which indicates that a byte of data has been received and is ready to be read from the SPART to the Processor. When the read operation is performed, RDA is reset. Transmit Buffer Ready (TBR) is a status signal which indicates that the transmit buffer in the SPART is ready to accept a byte for transmission. When a write operation is performed and the SPART is not ready for more transmission data, TBR is reset. The SPART is fully synchronous with the clock signal CLK; this implies that transfers between the driver and SPART can be controlled by applying IOCS, IOR/W, IOADDR, and DATABUS (in the case of a write operation) for a single clock cycle and capturing the transferred data on the next positive clock edge. The received data on RxD, however, is asynchronous with respect to CLK. Also, the serial I/O port on the workstation which receives the transmitted data from TxD has no access to CLK. This interface thus constitutes the “A” for “Asynchronous” in SPART and requires an understanding of RS-232 signal timing and (re)synchronization.

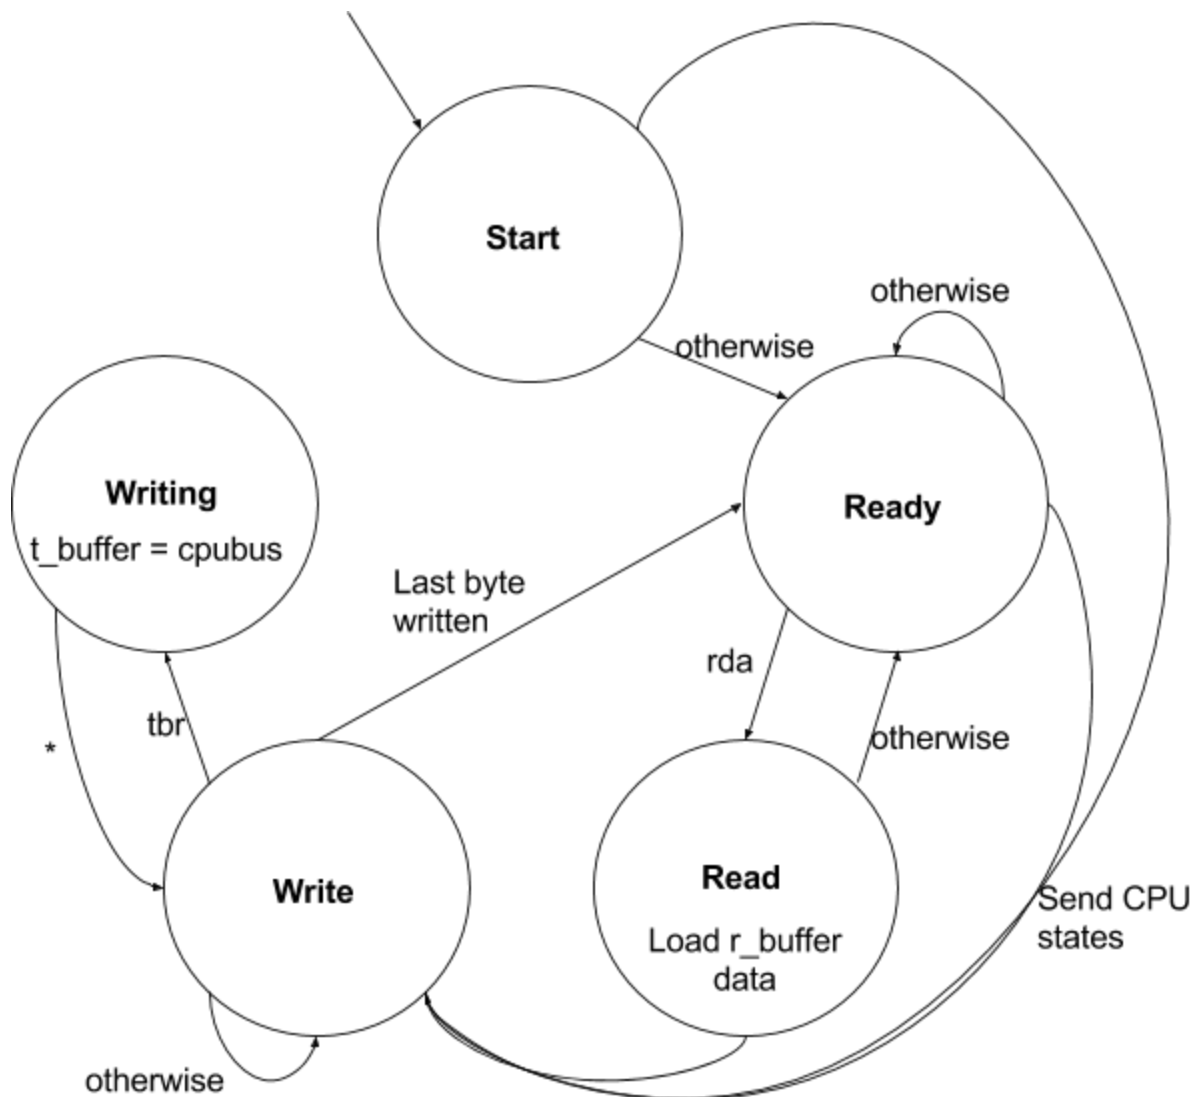
SPART Diagram & Interface



Controller Driver

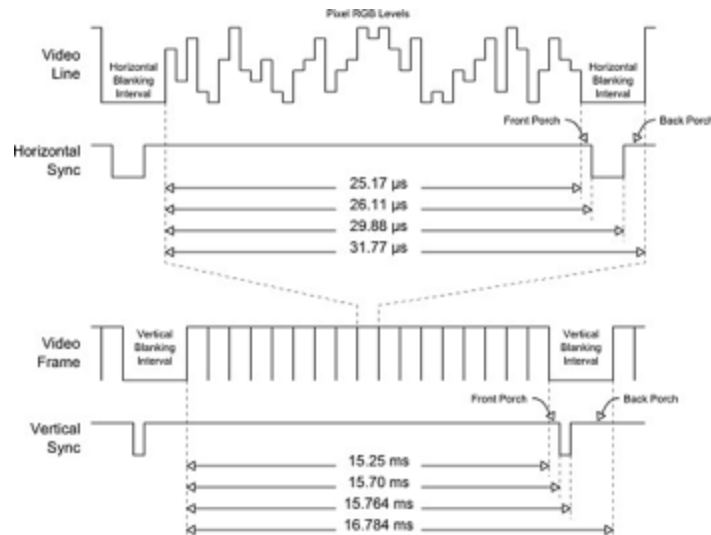
The controller driver is tasked with reloading the controller button states from the SPART receiver buffer when address \$4016 (or \$0 from controller's point of view) is set. In addition, the driver must grab the CPU databus on a read and place a button value on bit 0. On the first read, the button state of value A is placed on the databus, followed by B, Select, Start, Up, Down, Left, Right. The value will be 1 for pressed and 0 for not pressed. After reading the first 8 buttons, the driver will output a 0 on the databus. Lastly, the controller driver can also be used to control the SPART module to output to the UART port of the computer.

Controller Driver State Machine



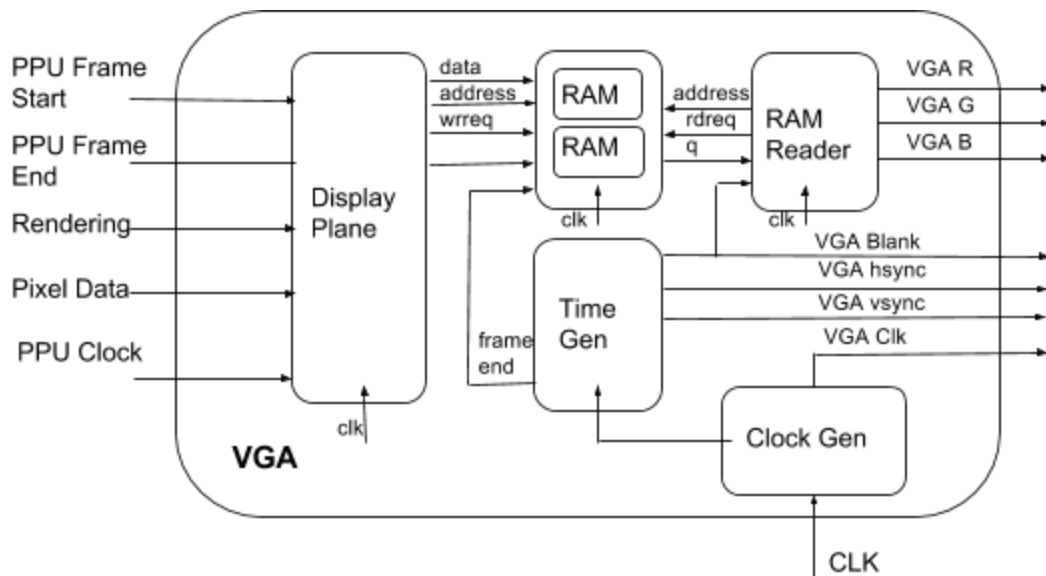
8. VGA

The VGA interface consists of sending the pixel data to the screen one row at a time from left to right. In between each row it requires a special signal called horizontal sync (hsync) to be asserted at a specific time when only black pixels are being sent, called the blanking interval. This happens until the bottom of the screen is reached when another blanking interval begins where the interface is only sending black pixels, but instead of hsync being asserted the vertical sync signal is asserted.



The main difficulty with the VGA interface will be designing a system to take the PPU output (a 256x240 image) and converting it into a native resolution of 640x480 or 1280x960. This was done by adding two RAM blocks to buffer the data.

VGA Diagram



VGA Interface

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
V_BLANK_N	output		Syncing each pixel
VGA_R[7:0]	output		Red pixel value
VGA_G[7:0]	output		Green pixel value
VGA_B[7:0]	output		Blue pixel value
VGA_CLK	output		VGA clock
VGA_HS	output		Horizontal line sync
VGA_SYNC_N	output		0
VGA_VS	output		Vertical line sync
pixel_data[7:0]	input	PPU	Pixel data to be sent to the display
ppu_clock	input	PPU	pixel data is updated every ppu clock cycle
rendering	input	PPU	high when PPU is rendering
frame_end	input	PPU	high at the end of a PPU frame
frame_start	input	PPU	high at start of PPU frame

VGA Clock Gen

This module takes in a 50MHz system clock and creates a 25.175MHz clock, which is the standard VGA clock speed.

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
VGA_CLK	output	VGA	Clock synced to VGA timing
locked	output		Locks VGA until clock is ready

VGA Timing Gen

This block is responsible for generating the timing signals for VGA with a screen resolution of 480x640. This includes the horizontal and vertical sync signals as well as the blank signal for each pixel.

Signal name	Signal Type	Source/Dest	Description
VGA_CLK	input	Clock Gen	vga_clk
rst_n	input		System active low reset
V_BLANK_N	output	VGA, Ram Reader	Syncing each pixel
VGA_HS	output	VGA	Horizontal line sync
VGA_VS	output	VGA	Vertical line sync

VGA Display Plane

The PPU will output sprite and background pixels to the VGA module, as well as enables for each. The display plane will update the RAM block at the appropriate address with the pixel data on every PPU clock cycle when the PPU is rendering.

Signal name	Signal Type	Source/Dest	Description
clk	input		System clock
rst_n	input		System active low reset
ppu_clock	input	PPU	Clock speed that the pixels from the PPU come in
wr_address	input	RAM	Address to write to
wr_req	output	RAM	Write data to the RAM
data_out[7:0]	output	RAM	The pixel data to store in RAM
pixel_data[7:0]	input	PPU	Pixel data to be sent to the display
rendering	input	PPU	high when PPU is rendering
frame_start	input	PPU	high at start of PPU frame

VGA RAM Wrapper

This module instantiates two 2-port RAM blocks and using control signals, it will have the PPU write to a specific RAM block, while the VGA reads from another RAM block. The goal of this module was to make sure that the PPU writes never overlap the VGA reads, because the PPU runs at a faster clock rate.

Signal name	Signal Type	Source/Dest	Description
clk	input		
rst_n	input		System active low reset
wr_address	input	Display Plane	Address to write to
wr_req	input	Display Plane	Request to write data
data_in[5:0]	input	Display Plane	The data into the RAM
rd_req	input	RAM Reader	Read data out from RAM
rd_address	input	RAM Reader	Address to read from
data_out[5:0]	output	RAM Reader	data out from RAM
ppu_frame_end	input	PPU	high at the end of a PPU frame
vga_frame_end	input	VGA	high at end of VGA frame

VGA RAM Reader

The RAM Reader is responsible for reading data from the correct address in the RAM block and outputting it as an RGB signal to the VGA. It will update the RGB signals every time the blank signal goes high. The NES supported a 256x240 image, which we will be converting to a 640x480 image. This means that the 256x240 image will be multiplied by 2, resulting in a 512x480 image. The remaining 128 pixels on the horizontal line will be filled with black pixels by this block. Lastly, this block will take use the pixel data from the PPU and the NES Palette RGB colors, to output the correct colors to the VGA.

Signal name	Signal Type	Source/Dest	Description
-------------	-------------	-------------	-------------

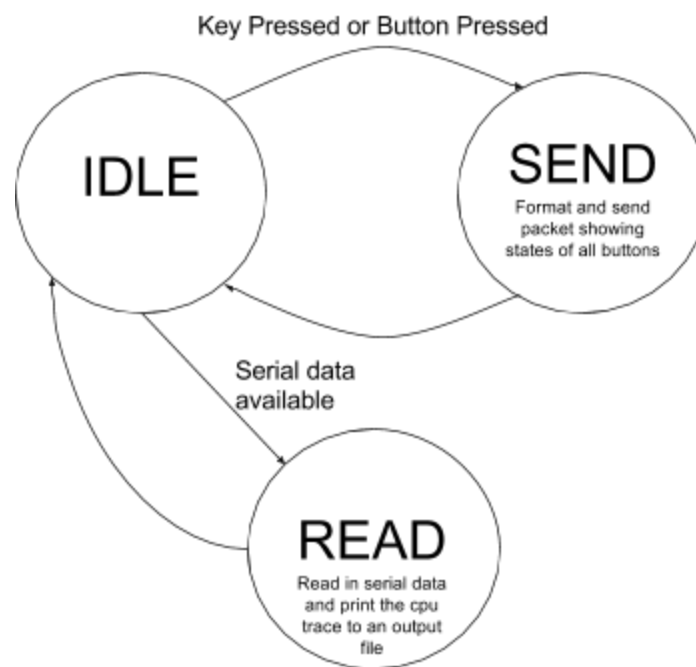
clk	input		
rst_n	input		System active low reset
rd_req	output	RAM	Read data out from RAM
rd_address	output	RAM	Address to read from
data_out[7:0]	input	RAM	data out from RAM
VGA_R[7:0]	output		VGA Red pixel value
VGA_G[7:0]	output		VGA Green pixel value
VGA_B[7:0]	output		VGA Blue pixel value
VGA_Blnk[7:0]	input	Time Gen	VGA Blank signal (high when we write each new pixel)

9. Software

Controller Simulator

In order to play games on the NES and provide input to our FPGA, we will have a java program that uses the JSSC (Java Simple Serial Connector) library to read and write data serially using the SPART interface. The program will provides a GUI that was created using the JFrame library. This GUI will respond to mouse clicks as well as key presses when the window is in focus. When a button state on the simulator is changed, it will trigger the program to send serial data. When data is detected on the rx line, the simulator will read in the data (every time there is 8 bytes in the buffer) and will output this data as a CPU trace to an output file. Instructions to invoke this program can be found in the README file of our github directory.

Controller Simulator State Machine



Controller Simulator Output Packet Format

Packet name	Packet type	Packet Format	Description
Controller Data	output	ABST-UDLR	This packet indicates which buttons are being pressed. A 1 indicates pressed, a 0 indicates not pressed. (A) A button, (B) B button, (S) Select button, (T) Start button, (U) Up, (D) Down, (L) Left, (R) Right

Controller Simulator GUI and Button Map

The NES controller had a total of 8 buttons, as shown below.



The NES buttons will be mapped to specific keys on the keyboard. The keyboard information will be obtained using KeyListeners in the java.awt.* library. The following table indicates how the buttons are mapped and their function in Super Mario Bros.

Keyboard button	NES Equivalent	Super Mario Bros. Function
X Key	A Button	Jump (Hold to jump higher)
Z Key	B Button	Sprint (Hold and use arrow keys)
Tab Key	Select Button	Pause Game
Enter Key	Start Button	Start Game
Up Arrow	Up on D-Pad	No function
Down Arrow	Down on D-Pad	Enter pipe (only works on some pipes)
Left Arrow	Left on D-Pad	Move left
Right Arrow	Right on D-Pad	Move right

Assembler

We will include an assembler that allows custom software to be developed for our console. This assembler will convert assembly code to machine code for the NES on .mif files that we can load into our FPGA. It will include support for labels and commenting. The ISA is specified in the table below:

Opcode Table

Opcode	Mode	Hex	Opcode	Mode	Hex	Opcode	Mode	Hex
ADC	Immediate	69	DEC	Zero Page	C6	ORA	Absolute	0D
ADC	Zero Page	65	DEC	Zero Page, X	D6	ORA	Absolute, X	1D
ADC	Zero Page, X	75	DEC	Absolute	CE	ORA	Absolute, Y	19
ADC	Absolute	6D	DEC	Absolute, X	DE	ORA	Indirect, X	01
ADC	Absolute, X	7D	DEX	Implied	CA	ORA	Indirect, Y	11
ADC	Absolute, Y	79	DEY	Implied	88	PHA	Implied	48
ADC	Indirect, X	61	EOR	Immediate	49	PHP	Implied	08
ADC	Indirect, Y	71	EOR	Zero Page	45	PLA	Implied	68
AND	Immediate	29	EOR	Zero Page, X	55	PLP	Implied	28
AND	Zero Page	25	EOR	Absolute	4D	ROL	Accumulator	2A
AND	Zero Page, X	35	EOR	Absolute, X	5D	ROL	Zero Page	26
AND	Absolute	2D	EOR	Absolute, Y	59	ROL	Zero Page, X	36
AND	Absolute, X	3D	EOR	Indirect, X	41	ROL	Absolute	2E
AND	Absolute, Y	39	EOR	Indirect, Y	51	ROL	Absolute, X	3E
AND	Indirect, X	21	INC	Zero Page	E6	ROR	Accumulator	6A
AND	Indirect, Y	31	INC	Zero Page, X	F6	ROR	Zero Page	66
ASL	Accumulator	0A	INC	Absolute	EE	ROR	Zero Page, X	76
ASL	Zero Page	06	INC	Absolute, X	FE	ROR	Absolute	6E
ASL	Zero Page, X	16	INX	Implied	E8	ROR	Absolute, X	7E
ASL	Absolute	0E	INY	Implied	C8	RTI	Implied	40
ASL	Absolute, X	1E	JMP	Indirect	6C	RTS	Implied	60
BCC	Relative	90	JMP	Absolute	4C	SBC	Immediate	E9
BCS	Relative	B0	JSR	Absolute	20	SBC	Zero Page	E5
BEQ	Relative	F0	LDA	Immediate	A9	SBC	Zero Page, X	F5

BIT	Zero Page	24	LDA	Zero Page	A5	SBC	Absolute	ED
BIT	Absolute	2C	LDA	Zero Page, X	B5	SBC	Absolute, X	FD
BMI	Relative	30	LDA	Absolute	AD	SBC	Absolute, Y	F9
BNE	Relative	D0	LDA	Absolute, X	BD	SBC	Indirect, X	E1
BPL	Relative	10	LDA	Absolute, Y	B9	SBC	Indirect, Y	F1
BRK	Implied	00	LDA	Indirect, X	A1	SEC	Implied	38
BVC	Relative	50	LDA	Indirect, Y	B1	SED	Implied	F8
BVS	Relative	70	LDX	Immediate	A2	SEI	Implied	78
CLC	Implied	18	LDX	Zero Page	A6	STA	Zero Page	85
CLD	Implied	D8	LDX	Zero Page, Y	B6	STA	Zero Page, X	95
CLI	Implied	58	LDX	Absolute	AE	STA	Absolute	8D
CLV	Implied	B8	LDX	Absolute, Y	BE	STA	Absolute, X	9D
CMP	Immediate	C9	LDY	Immediate	A0	STA	Absolute, Y	99
CMP	Zero Page	C5	LDY	Zero Page	A4	STA	Indirect, X	81
CMP	Zero Page, X	D5	LDY	Zero Page, X	B4	STA	Indirect, Y	91
CMP	Absolute	CD	LDY	Absolute	AC	STX	Zero Page	86
CMP	Absolute, X	DD	LDY	Absolute, X	BC	STX	Zero Page, Y	96
CMP	Absolute, Y	D9	LSR	Accumulator	4A	STX	Absolute	8E
CMP	Indirect, X	C1	LSR	Zero Page	46	STY	Zero Page	84
CMP	Indirect, Y	D1	LSR	Zero Page, X	56	STY	Zero Page, X	94
CPX	Immediate	E0	LSR	Absolute	4E	STY	Absolute	8C
CPX	Zero Page	E4	LSR	Absolute, X	5E	TAX	Implied	AA
CPX	Absolute	EC	NOP	Implied	EA	TAY	Implied	A8
CPY	Immediate	C0	ORA	Immediate	09	TSX	Implied	BA
CPY	Zero Page	C4	ORA	Zero Page	05	TXA	Implied	8A
CPY	Absolute	CC	ORA	Zero Page, X	15	TXS	Implied	9A
						TYA	Implied	98

NES Assembly Formats

Our assembler will allow the following input format, each instruction/label will be on its own line. In addition unlimited whitespace is allowed:

Instruction Formats		
Instruction Type	Format	Description
Constant	Constant_Name = <Constant Value>	Must be declared before CPU_Start
Label	Label_Name:	Cannot be the same as an opcode name. Allows reference from branch opcodes.
Comment	; Comment goes here	Anything after the ; will be ignored
CPU Start	_CPU:	Signals the start of CPU memory
Accumulator	<OPCODE>	Accumulator is value affected by Opcode
Implied	<OPCODE>	Operands implied by opcode. ie. TXA has X as source and Accumulator as destination
Immediate	<OPCODE> #<Immediate>	The decimal number will be converted to binary and used as operand
Absolute	<OPCODE> \$<ADDR/LABEL>	The byte at the specified address is used as operand
Zero Page	<OPCODE> \$<BYTE OFFSET>	The byte at address \$00XX is used as operand.
Relative	<OPCODE> \$<BYTE OFFSET/LABEL>	The byte at address PC +/- Offset is used as operand. Offset can range -128 to +127
Absolute Index	<OPCODE> \$<ADDR/LABEL>,<X or Y>	Absolute but value in register added to address.
Zero Page Index	<OPCODE> \$<BYTE OFFSET>,<X or Y>	Zero page but value in register added to offset.
Zero Page X Indexed Indirect	<OPCODE> (\$<BYTE OFFSET>,<X>)	Value in X added to offset. Address in \$00XX (where XX is new offset) is used as the address for the operand.
Zero Page Y Indexed Indirect	<OPCODE> (\$<BYTE OFFSET>,<Y>)	The address in \$00XX, where XX is byte offset, is added to the value in Y and is used as the address for the operand.
Data instruction	<.DB or .DW> <data values>	If .db then the data values must be bytes, if .dw then the data values must be 2 bytes. Multiple comma separated data values can be include for each instruction. Constants are valid.

Number Formats		
Immediate Decimal (Signed)	#<(-)DDD>	Max 127, Min -128
Immediate Hexadecimal (Signed)	#\$<HH>	
Immediate Binary (Signed)	##<BBBB.BBBB>	Allows ' ' in between bits
Address/Offset Hex	\$<Addr/Offset>	8 bits offset, 16 bits address
Address/Offset Binary	\$%<Addr/Offset>	8 bits offset, 16 bits address
Offset Decimal (Relative only)	#<(-)DDD>	Relative instructions can't be Immediate, so this is allowed. Max 127, Min -128
Constant first byte	<Constant_Name	
Constant second byte	>Constant_Name	
Constant	Constant_Name	
Label	Label_Name	Not valid for data instructions

Invoking Assembler

Usage	Description
java NESAssemble <input file> <cpuoutput.mif> <ppuoutput.mif>	Reads the input file and outputs the CPU ROM to cpuoutput.mif and the PPU ROM to ppuoutput.mif

iNES ROM Converter

Most NES games are currently available online in files of the iNES format, a header format used by most software NES emulators. Our NES will not support this file format. Instead, we will write a java program that takes an iNES file as input and outputs two .mif files that contain the CPU RAM and the PPU VRAM. These files will be used to instantiate the ROM's of the CPU and PPU in our FPGA.

Usage	Description
java NESToMIF <input.nes> <cpuoutput.mif> <ppuoutput.mif>	Reads the input file and outputs the CPU RAM to cpuoutput.mif and the PPU VRAM to ppuoutput.mif

Tic Tac Toe

We also implemented Tic Tac Toe in assembly for initial integration tests. We bundled it into a NES ROM, and thus can run it on existing emulators as well as our own hardware.

10. Testing & Debug

Our debugging process had multiple steps

Simulation

For basic sanity check, we simulated each module independently to make sure the signals behave as expected.

Test

For detailed check, we wrote an automated testbench to confirm the functionality. The CPU test suite was from <https://github.com/Klaus2m5/> and we modified the test suite to run on the fceux NES emulator.

Integrated Simulation

After integration, we simulated the whole system with the ROM installed. We were able to get a detailed information at each cycle but the simulation took too long. It took about 30 minutes to simulate CPU operation for one second. Thus we designed a debug and trace module in hardware that could output CPU traces during actual gameplay..

Tracer

We added additional code in Controller so that the Controller can store information at every cycle and dump them back to the serial console under some condition. At first, we used a button as the trigger, but after we analyzed the exact problem, we used a conditional statement(for e.g. when PC reached a certain address) to trigger the dump. When the condition was met, the Controller would stall the CPU and start dumping what the CPU has been doing, in the opposite order of execution. The technique was extremely useful because we came to the conclusion that there must be a design defect when Mario crashed, such as using don't cares or high impedance. After we corrected the defect, we were able to run Mario.

11. Results

We were able to get NES working, thanks to our rigorous verification process, and onboard debug methodology. Some of the games we got working include Super Mario Bros, Galaga, Tennis, Golf, Donkey Kong, Ms Pacman, Defender II, Pinball, and Othello.

12. Possible Improvements

- Create a working audio processing unit
- More advanced memory mapper support
- Better image upscaling such as hqx
- Support for actual NES game carts
- HDMI
- VGA buffer instead of two RAM blocks to save space

13. References and Links

Ferguson, Scott. "PPU Tutorial." N.p., n.d. Web. <<https://opcode-defined.quora.com>>.

"6502 Specification." NesDev. N.p., n.d. Web. 10 May 2017. <<http://nesdev.com/6502.txt>>.

Dormann, Klaus. "6502 Functional Test Suite." *GitHub*. N.p., n.d. Web. 10 May 2017. <<https://github.com/Klaus2m5/>>.

"NES Reference Guide." Nesdev. N.p., n.d. Web. 10 May 2017. <http://wiki.nesdev.com/w/index.php/NES_reference_guide>.

Java Simple Serial Connector library:
<<https://github.com/scream3r/java-simple-serial-connector>>

Final Github release: <https://github.com/jtgebert/fpganes_release>

14. Contributions

Eric Sullivan

Designed and debugged the NES picture processing unit, created a comprehensive set of PPU testbenches to verify functionality, Integrated the VGA to the PPU, implemented the DMA and dummy APU, started a CPU simulator in python, Helped debug the integrated system.

Patrick Yang

Specified the CPU microarchitecture along with Pavan, designed the ALU, registers, and memory interface unit, wrote a self checking testbench, responsible for CPU debug, integrated all modules on top level file, and debugged of the integrated system. Helped Jon to modify controller driver to also be an onboard CPU trace module.

Pavan Holla

Specified the CPU microarchitecture along with Patrick, designed the decoder and interrupt handler, and wrote the script that generates the processor control module. Modified a testsuite and provided the infrastructure for CPU verification. Wrote tic tac toe in assembly as a fail-safe game. Also, worked on a parallel effort to integrate undocumented third party NES IP.

Jonathan Ebert

Modified the VGA to interface with the PPU. Wrote a new driver to control our existing SPART module to act as a NES controller and as an onboard CPU trace module. Wrote Java program to communicate with the SPART module using the JSSC library. Wrote memory wrappers, hardware decoder, and generated all game ROMs. Helped debug the integrated system. Wrote a very simple assembler. Wrote script to convert NES ROMs to MIF files. Also, worked on a parallel effort to integrate undocumented third party NES IP.